


Deep Learning Tutorial

李宏毅

Hung-yi Lee

Deep learning attracts lots of attention.

- I believe you have seen lots of exciting results before.



Deep learning trends at Google. Source: SIGMOD/Jeff Dean

This talk focuses on the basic techniques.

Outline

Lecture I: Introduction of Deep Learning



Lecture II: Tips for Training Deep Neural Network



Lecture III: Variants of Neural Network



Lecture IV: Next Wave

Lecture I: Introduction of Deep Learning

Three Steps for Deep Learning



Three Steps for Deep Learning



based on
training data

Three Steps for Deep Learning

- Speech Recognition

$$f^*(\text{sound waveform}) = \text{“你好”}$$

- Handwritten Recognition

$$f^*(\text{handwritten digit}) = \text{“2”}$$


- Playing Go

$$f^*(\text{Go board state}) = \text{“5-5” (step)}$$

- Dialogue System

$$f^*(\text{“Hi” (what the user said)}) = \text{“Hello” (system response)}$$

Step 3:
Learn!




Pick the
best
function f^*


Three Steps for Deep Learning




Deep Learning is so simple



Three Steps for Deep Learning




Human Brains



Neural Network


Neuron

$$z = a_1 w_1 + \dots + a_k w_k + \dots + a_K w_K + b$$




Neural Network

Neuron




Neural Network

Different connections leads to different network structured




Each neurons can have different values of weights and biases.

Weights and biases are network parameters θ




Fully Connect Feedforward Network




Sigmoid Function


$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Fully Connect Feedforward Network



Fully Connect Feedforward Network




Network is a function.
Input vector, output vector

$$f \left(\begin{bmatrix} 1 \\ -1 \end{bmatrix} \right) = \begin{bmatrix} 0.62 \\ 0.83 \end{bmatrix} \quad f \left(\begin{bmatrix} 0 \\ 0 \end{bmatrix} \right) = \begin{bmatrix} 0.51 \\ 0.85 \end{bmatrix}$$

Given parameters θ , define a function

Given network structure, define a function set

Fully Connect Feedforward Network




Deep means many hidden layers

Ultra Deep Network

http://cs231n.stanford.edu/slides/winter1516_lecuture8.pdf

16.4%



AlexNet (2012)

8 layers


7.3%



VGG (2014)

19 layers

6.7%



GoogleNet (2014)

22 layers

Ultra Deep Network

This ultra deep network
have special structure.

(Lecture IV)

3.57%

16.4%



AlexNet
(2012)

7.3%



VGG
(2014)

6.7%




GoogleNet
(2014)

Residual Net
(2015)




101 layers




Taipei
101

Fully Connect Feedforward Network - Matrix Operation



Fully Connect Feedforward Network - Matrix Operation



$$y = f(x)$$

Using parallel computing techniques (e.g. GPU)
to speed up matrix operation

$$= \sigma(W^L \cdots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \cdots + b^L)$$

Output Layer (Option)

- Softmax layer as the output layer

Ordinary Layer

$$z_1 \rightarrow \sigma \rightarrow y_1 = \sigma(z_1)$$

$$z_2 \rightarrow \sigma \rightarrow y_2 = \sigma(z_2)$$

$$z_3 \rightarrow \sigma \rightarrow y_3 = \sigma(z_3)$$

In general, the output of network can be any value.

May not be easy to interpret


Output Layer (Option)

- Softmax layer as the output layer


Probability:

- $1 > y_i > 0$
- $\sum_i y_i = 1$

Softmax Layer




FAQ



- Q: How many layers? How many neurons for each layer?
- Q: Can the structure be automatically determined?


Three Steps for Deep Learning



Example Application




Input



Ink → 1
No ink → 0


Output




Each dimension represents the confidence of a digit.

Example Application


- Handwriting Digit Recognition



Fully Connect Feedforward Network



Fully Connect Feedforward Network



Step 2 Define the goodness of function based on training data

Step 3 Pick the best function

Training Data

- Preparing training data: images and their labels




“5”



“0”



“4”




“1”



“9”



“2”




“1”



“3”

The learning target is defined on
the training data.

Learning Target



Ink \rightarrow 1

No ink \rightarrow 0


The learning target is

Input: $\rightarrow y_1$ has the maximum value

Input: $\rightarrow y_2$ has the maximum value

LOSS

A good function should make the loss of all examples as small as possible.




Loss can be **square error** or **cross entropy**
between the network output and target

target

Total Loss

For all training data ...



Total Loss:


$$L = \sum_{r=1}^R l_r$$

As small as possible

Find a function in function set that minimize total loss L

Find the network parameters θ^* that minimize total loss L

Three Steps for Deep Learning



How to pick the best function


Find network parameters θ^* that minimize total loss L

Enumerate all possible values

Network parameters $\theta =$
 $\{w_1, w_2, w_3, \dots, b_1, b_2, b_3, \dots\}$

Millions of parameters


E.g. speech recognition: 8 layers and
1000 neurons each layer



Gradient Descent

Network parameters $\theta = \{w_1, w_2, \dots, b_1, b_2, \dots\}$


Find network parameters θ^* that minimize total loss L



Gradient Descent

Network parameters $\theta = \{w_1, w_2, \dots, b_1, b_2, \dots\}$


Find network parameters θ^* that minimize total loss L



Gradient Descent

Network parameters $\theta = \{w_1, w_2, \dots, b_1, b_2, \dots\}$


Find network parameters θ^* that minimize total loss L




Gradient Descent

Network parameters $\theta = \{w_1, w_2, \dots, b_1, b_2, \dots\}$

Find network parameters θ^* that minimize total loss L




Gradient Descent




$$\nabla L = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \frac{\partial L}{\partial w_2} \\ \vdots \\ \frac{\partial L}{\partial b_1} \\ \vdots \end{bmatrix}$$

gradient

Gradient Descent




Gradient Descent



Gradient Descent

Hopfully, we would reach
a minima




Gradient Descent

- When considering multiple parameters together, do you see any problem?

Local Minima



Who is Afraid of Non-Convex
Loss Functions?
[http://videolectures.net/eml07
_lecun_wia/](http://videolectures.net/eml07_lecun_wia/)

- Gradient descent never guarantee global minima



想像你在玩世紀帝國.....

沒有探索過的地方被戰霧覆蓋




Gradient Descent

This is the “learning” of machines in deep learning

→ Even alpha go using this approach.

大家以為 Learning 是




其實 Learning 只是



I hope you are not too disappointed :p

Backpropagation

- Backpropagation: an efficient way to compute $\partial L / \partial w$
 - Ref:
http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS_2015_2/Lecture/DNN%20backprop.ecm.mp4/index.html



theano

libdnn
台大周伯威
同學開發

Caffe


Microsoft
CNTK



mxnet


Don't worry about $\partial L / \partial w$, the toolkits will handle it.

You can do lots of different things




For example, you can do


- Image Recognition



For example, you can do



Document Classification



Playing Go




19 x 19 image
(matrix)

Black: 1
white: -1
none: 0



Network



Next move
(19 x 19
positions)


19 x 19 vector

Fully-connected feedword
network can be used



But CNN performs much better.

(Lecture III)

Playing Go




Training:



Concluding Remarks

- Deep Learning is simple & powerful!

但 Deep Learning 就像 雷神之槌



無法輕易被舉起來

<http://ent.ltn.com.tw/news/breakingnews/1144545>

Lecture II: Tips for Training DNN

Outline of Lecture II

“Hello World” for Deep Learning


Recipe of Deep Learning

Keras

If you want to learn theano:

http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS_2015_2/Lecture/Theano%20DNN.ecm.mp4/index.html

[http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS_2015_2/Lecture/RNN%20training%20\(v6\).ecm.mp4/index.html](http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS_2015_2/Lecture/RNN%20training%20(v6).ecm.mp4/index.html)



Interface of
TensorFlow or
Theano

Very flexible
Need some
effort to learn


Easy to learn and use
(still have some flexibility)
You can modify it if you can write
TensorFlow or Theano

Keras

- François Chollet is the author of Keras.
 - He currently works for Google as a deep learning engineer and researcher.
- Keras means *horn* in Greek
- Documentation: <http://keras.io/>
- Example:
<https://github.com/fchollet/keras/tree/master/examples>

使用 Keras 心得


Deep Learning研究生




朋友覺得我在



我媽覺得我在



大眾覺得我在



指導教授覺得我在




我以為我在



事實上我在

Example Application

- Handwriting Digit Recognition



MNIST Data: <http://yann.lecun.com/exdb/mnist/>
“Hello world” for deep learning

Keras provides data sets loading function: <http://keras.io/datasets/>


Keras

Step 1:
Network
Structure

Step 2:
Learning
Target

Step 3:
Learn!

28x28




```
model = Sequential()
```

```
model.add( Dense( input_dim=28*28,  
                  output_dim=500 ) )  
model.add( Activation('sigmoid') )
```

```
model.add( Dense( output_dim=500 ) )  
model.add( Activation('sigmoid') )
```

```
model.add( Dense( output_dim=10 ) )  
model.add( Activation('softmax') )
```

Keras



```
model.compile(loss='mse',
               optimizer=SGD(lr=0.1),
               metrics=['accuracy'])
```

Keras



Step 3.1: Configuration

```
model.compile(loss='mse',  
               optimizer=SGD(lr=0.1),  
               metrics=['accuracy'])
```

$$w \leftarrow w - \eta \partial L / \partial w$$

0.1

Step 3.2: Find the optimal network parameters

```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```

Training data
(Images)

Labels
(digits)

等一下再講


Keras



Step 3.2: Find the optimal network parameters


```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```

numpy array




Number of training examples

numpy array



Number of training examples

Keras



Save and load models

<http://keras.io/getting-started/faq/#how-can-i-save-a-keras-model>

How to use the neural network (testing):

```
score = model.evaluate(x_test, y_test)
case 1: print('Total loss on Testing Set:', score[0])
          print('Accuracy of Testing Set:', score[1])
```

```
case 2: result = model.predict(x_test)
```

Keras


- Using GPU to speed training
 - Way 1
 - THEANO_FLAGS=device=gpu0 python YourCode.py
 - Way 2 (in your code)
 - import os
 - os.environ["THEANO_FLAGS"] = "device=cpu"

Outline of Lecture II



“Hello World” for Deep Learning

Recipe of Deep Learning

Recipe of Deep Learning



Do not always blame Overfitting




Deep Residual Learning for Image Recognition
<http://arxiv.org/abs/1512.03385>

Recipe of Deep Learning

Different approaches for different problems.

e.g. dropout for good results on testing data



Neural Network

Good Results on Testing Data?



Good Results on Training Data?

Other methods do not emphasize this.




YES

Recipe of Deep Learning



Choosing Proper Loss



Square
Error

$$\sum_{i=1}^{10} (y_i - \hat{y}_i)^2 = 0$$


Cross
Entropy

$$-\sum_{i=1}^{10} \hat{y}_i \ln y_i = 0$$

target

Choosing Proper Loss

When using softmax output layer,
choose cross entropy



Let's try it

Square Error

```
model.compile(loss='mse',  
              optimizer=SGD(lr=0.1),  
              metrics=['accuracy'])
```


Cross Entropy

```
model.compile(loss='categorical_crossentropy',  
              optimizer=SGD(lr=0.1),  
              metrics=['accuracy'])
```


Let's try it

Testing:

	Accuracy
Square Error	0.11
Cross Entropy	0.84



Recipe of Deep Learning




```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```


We do not really minimize total loss!

Mini-batch

Mini-batch



Mini-batch



- Randomly initialize network parameters

- Pick the 1st batch
 $L' = l^1 + l^{31} + \dots$
Update parameters once
- Pick the 2nd batch
 $L'' = l^2 + l^{16} + \dots$
Update parameters once
- ⋮
- Until all mini-batches have been picked


one epoch

Repeat the above process

Mini-batch

```
model.fit(x_train, y_train, batch_size=100, nb_epoch=20)
```

Mini-batch



100 examples in a mini-batch


Repeat 20 times

- Pick the 1st batch
 $L' = l^1 + l^{31} + \dots$
Update parameters once
 - Pick the 2nd batch
 $L'' = l^2 + l^{16} + \dots$
Update parameters once
⋮
 - Until all mini-batches have been picked
- one epoch


We do not really minimize total loss!

Mini-batch

Mini-batch



Mini-batch




- Randomly initialize network parameters
- Pick the 1st batch
 $L' = l^1 + l^{31} + \dots$
Update parameters once
- Pick the 2nd batch
 $L'' = l^2 + l^{16} + \dots$
Update parameters once
⋮

L is different each time when we update parameters!


目標換來換去?!

Mini-batch

Original Gradient Descent



With Mini-batch




The colors represent the total loss.

Mini-batch is Faster

Not always true with parallel computing.

Original Gradient Descent


Update after seeing all examples



Mini-batch has better performance!

With Mini-batch


If there are 20 batches, update 20 times in one epoch.




Testing:

Mini-batch is Better!


	Accuracy
Mini-batch	0.84
No batch	0.12




Shuffle the training examples for each epoch



Recipe of Deep Learning




Hard to get the power of Deep ...




Let's try it

Testing:

	Accuracy
3 layers	0.84
9 layers	0.11



Vanishing Gradient Problem



Smaller gradients

Learn very slow

Almost random

Larger gradients



Learn very fast

Already converge

based on random!?

Vanishing Gradient Problem


Smaller gradients



Intuitive way to compute the derivatives ...


$$\frac{\partial l}{\partial w} = ? \quad \frac{\Delta l}{\Delta w}$$

Hard to get the power of Deep ...



ReLU

- Rectified Linear Unit (ReLU)




[Xavier Glorot, AISTATS'11]
[Andrew L. Maas, ICML'13]
[Kaiming He, arXiv'15]

Reason:


1. Fast to compute
2. Biological reason
3. Infinite sigmoid with different biases
4. Vanishing gradient problem

ReLU




ReLU

A Thinner linear network



Let's try it

```
model.add( Activation('sigmoid') )
```




```
model.add( Activation('relu') )
```

Let's try it

- 9 layers


Testing:

9 layers	Accuracy
Sigmoid	0.11
ReLU	0.96




ReLU - variant

Leaky ReLU



Parametric ReLU




α also learned by
gradient descent

Maxout

ReLU is a special cases of Maxout

- Learnable activation function [Ian J. Goodfellow, ICML'13]




You can have more than 2 elements in a group.

Maxout


ReLU is a special cases of Maxout

- Learnable activation function [Ian J. Goodfellow, ICML'13]
 - Activation function in maxout network can be any piecewise linear convex function
 - How many pieces depending on how many elements in a group



2 elements in a group



3 elements in a group




Recipe of Deep Learning




Learning Rates

Set the learning rate η carefully



Learning Rates

Set the learning rate η carefully



Learning Rates

- Popular & Simple Idea: Reduce the learning rate by some factor every few epochs.
 - At the beginning, we are far from the destination, so we use larger learning rate
 - After several epochs, we are close to the destination, so we reduce the learning rate
 - E.g. 1/t decay: $\eta^t = \eta / \sqrt{t + 1}$
- Learning rate cannot be one-size-fits-all
 - Giving different parameters different learning rates

Adagrad

Original: $w \leftarrow w - \eta \partial L / \partial w$

Adagrad: $w \leftarrow w - \boxed{\eta_w} \partial L / \partial w$

Parameter dependent learning rate

$$\eta_w = \frac{\eta}{\sqrt{\sum_{i=0}^t (g^i)^2}}$$

constant

g^i is $\partial L / \partial w$ obtained at the i-th update

Summation of the square of the previous derivatives

Adagrad

$$\eta_w = \frac{\eta}{\sqrt{\sum_{i=0}^t (g^i)^2}}$$

w_1	\mathbf{g}^0
	0.1

Learning rate:

$$\frac{\eta}{\sqrt{0.1^2}}$$

$$= \frac{\eta}{0.1}$$



$$\frac{\eta}{\sqrt{0.1^2 + 0.2^2}}$$

$$= \frac{\eta}{0.22}$$

w_2	\mathbf{g}^0
	20.0

Learning rate:

$$\frac{\eta}{\sqrt{20^2}}$$

$$= \frac{\eta}{20}$$




$$\frac{\eta}{\sqrt{20^2 + 10^2}} = \frac{\eta}{22}$$

- Observation:**
1. Learning rate is smaller and smaller for all parameters
 2. Smaller derivatives, larger learning rate, and vice versa

Why?

Larger derivatives

Smaller Learning Rate



Smaller Derivatives

Larger Learning Rate





2. Smaller derivatives, larger learning rate, and vice versa

Why?


Not the whole story

- Adagrad [John Duchi, JMLR'11]
- RMSprop
 - <https://www.youtube.com/watch?v=O3sxAc4hxZU>
- Adadelta [Matthew D. Zeiler, arXiv'12]
- “No more pesky learning rates” [Tom Schaul, arXiv'12]
- AdaSecant [Caglar Gulcehre, arXiv'14]
- Adam [Diederik P. Kingma, ICLR'15]
- Nadam
 - http://cs229.stanford.edu/proj2015/054_report.pdf

Recipe of Deep Learning




Hard to find optimal network parameters




In physical world

- Momentum



Momentum

Still not guarantee reaching global minima, but give some hope



Adam

RMSProp (Advanced Adagrad) + Momentum

```
model.compile(loss='categorical_crossentropy',
               optimizer=SGD(lr=0.1),
               metrics=['accuracy'])
```

```
model.compile(loss='categorical_crossentropy',
               optimizer=Adam(),
               metrics=['accuracy'])
```


Algorithm 1: *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation. g_t^2 indicates the elementwise square $g_t \odot g_t$. Good default settings for the tested machine learning problems are $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$. All operations on vectors are element-wise. With β_1^t and β_2^t we denote β_1 and β_2 to the power t .

Require: α : Stepsize
Require: $\beta_1, \beta_2 \in [0, 1]$: Exponential decay rates for the moment estimates
Require: $f(\theta)$: Stochastic objective function with parameters θ
Require: θ_0 : Initial parameter vector
 $m_0 \leftarrow 0$ (Initialize 1st moment vector)
 $v_0 \leftarrow 0$ (Initialize 2nd moment vector)
 $t \leftarrow 0$ (Initialize timestep)
while θ_t not converged **do**
 $t \leftarrow t + 1$
 $g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)
 $m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)
 $v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)
 $\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)
 $\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)
 $\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)
end while
return θ_t (Resulting parameters)



Let's try it

- ReLU, 3 layer

	Accuracy
Original	0.96
Adam	0.97



Recipe of Deep Learning




Why Overfitting?

- Training data and testing data can be different.

Training Data:



Testing Data:



Learning target is defined by the training data.


The parameters achieving the learning target do not necessarily have good results on the testing data.

Panacea for Overfitting


- Have more training data
- *Create* more training data (?)

Handwriting recognition:

Original
Training Data:



Created
Training Data:



Shift 15 °

Why Overfitting?

- For experiments, we added some noises to the testing data

```
-1.36230370e-01, 1.03749340e-01, 1.15432226e-01,  
2.58670464e-01, 1.48774333e+00, 1.92885328e+00,  
1.70038673e+00, 2.46242981e+00, 1.21244572e+00,  
-9.28660713e-01, 3.63209761e-01, -1.81327713e+00,  
-1.97910760e-01, 4.32874592e-01, -5.40565788e-01,  
2.95630655e-01, 2.07984424e+00, -1.84243292e+00,  
-5.11166017e-01, -5.80935128e-01, 1.06273647e+00,  
1.80551097e-02, 2.27983997e-02, -1.67979148e+00,  
8.12423001e-01, -6.25888706e-01, -1.25027082e+00,  
6.15135458e-01, -1.21394611e-01, -1.28089527e+00,  
3.24609806e-01, 6.70569391e-01, 1.49161323e-01,  
8.01573609e-01, 6.43116741e-01, -9.37629233e-02,  
1.74677366e+00, 6.80996008e-01, -7.03717611e-01,  
1.02079749e-01, 1.19505614e+00, -2.77959386e-01,  
-5.21652916e-02, 3.53683601e-01, -4.08310762e-01,  
-1.81042967e+00, -9.03308062e-01, 1.05404509e+00,  
-9.80876877e-01, 3.52078891e-01, 6.65981840e-01,  
1.06550150e+00, -2.28433613e-01, 3.64483904e-01,  
-1.51484666e+00, -7.52612872e-02, -2.97058082e-01,  
-7.27414382e-01, -2.45875340e-01, -1.27948942e-01,  
-3.69310620e-01, -2.62300428e+00, 2.11585073e+00,  
6.85561585e-01, -1.57443985e-01, 1.38128777e+00,  
6.84265587e-02, 3.12536292e-01, 4.54253185e-01,  
-7.88471875e-01, -6.58403343e-02, -1.41847985e+00,  
-1.39753340e-01, -5.55354856e-01, -5.01917779e-01,  
6.93118522e-01, -2.45360497e-01, -1.26943186e+00,  
-2.62323855e-01])  
n [3]: x test[0]
```



Why Overfitting?

- For experiments, we added some noises to the testing data


Testing:	Accuracy
Clean	0.97
Noisy	0.50

Training is not influenced.

Recipe of Deep Learning





Early Stopping




Keras: <http://keras.io/getting-started/faq/#how-can-i-interrupt-training-when-the-validation-loss-isnt-decreasing-anymore>

Recipe of Deep Learning




Weight Decay

- Our brain prunes out the useless link between neurons.





Doing the same thing to machine's brain improves the performance.



Source: Rethinking the Brain, Families and Work Institute, Rima Shore, 1997; Founders Network slide

Weight Decay



Weight decay is one kind of regularization

Close to zero (萎縮了)

Weight Decay

- Implementation

$$\text{Original: } w \leftarrow w - \eta \frac{\partial L}{\partial w}$$

$$\lambda = 0.01$$


Weight Decay:

$$w \leftarrow \underbrace{0.99}_{\downarrow} w - \eta \frac{\partial L}{\partial w}$$

Smaller and smaller


Keras: <http://keras.io/regularizers/>

Recipe of Deep Learning



Dropout


Training:



- **Each time before updating the parameters**
 - Each neuron has $p\%$ to dropout

Dropout

Training:




- **Each time before updating the parameters**
 - Each neuron has $p\%$ to dropout
 - ➡ **The structure of the network is changed.**
 - Using the new network for training

For each mini-batch, we resample the dropout neurons

Dropout


Testing:



➤ No dropout

- If the dropout rate at training is $p\%$,
all the weights times $(1-p)\%$
- Assume that the dropout rate is 50%.
If a weight $w = 1$ by training, set $w = 0.5$ for testing.

Dropout - Intuitive Reason




- When teams up, if everyone expect the partner will do the work, nothing will be done finally.
- However, if you know your partner will dropout, you will do better.
- When testing, no one dropout actually, so obtaining good results eventually.

Dropout - Intuitive Reason

- Why the weights should multiply $(1-p)\%$ (dropout rate) when testing?


Training of Dropout

Assume dropout rate is 50%




Testing of Dropout

No dropout




Weights from training



Weights multiply $(1-p)\%$

Dropout is a kind of ensemble.


Ensemble




Train a bunch of networks with different structures

Dropout is a kind of ensemble.

Ensemble




Dropout is a kind of ensemble.



- Using one mini-batch to train one network
- Some parameters in the network are shared

Dropout is a kind of ensemble.


Testing of Dropout



More about dropout

- More reference for dropout [Nitish Srivastava, JMLR'14] [Pierre Baldi, NIPS'13][Geoffrey E. Hinton, arXiv'12]
- Dropout works better with Maxout [Ian J. Goodfellow, ICML'13]
- Dropconnect [Li Wan, ICML'13]
 - Dropout delete neurons
 - Dropconnect deletes the connection between neurons
- Annealed dropout [S.J. Rennie, SLT'14]
 - Dropout rate decreases by epochs
- Standout [J. Ba, NISP'13]
 - Each neural has different dropout rate

Let's try it



```
model = Sequential()
```

```
model.add( Dense( input_dim=28*28,  
                  output_dim=500 ) )  
model.add( Activation('sigmoid') )
```

model.add(dropout(0.8))

```
model.add( Dense( output_dim=500 ) )  
model.add( Activation('sigmoid') )
```


model.add(dropout(0.8))

```
model.add( Dense(output_dim=10) )  
model.add( Activation('softmax') )
```




YES

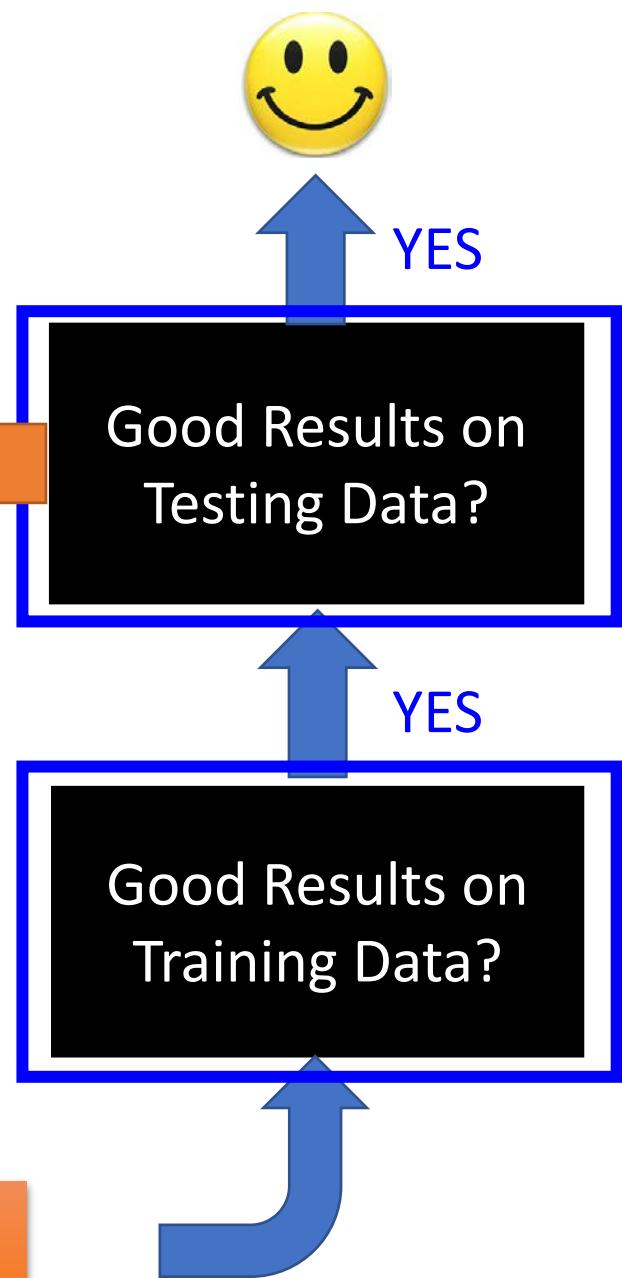

Let's try it



No Dropout




Recipe of Deep Learning



Concluding Remarks of Lecture II

Recipe of Deep Learning



Lecture III:

Variants of Neural Networks

Variants of Neural Networks


Convolutional Neural
Network (CNN)

Considering the
property of images

Recurrent Neural Network
(RNN)

Why CNN for Image?

- When processing image, the first layer of fully connected network would be very large




Can the fully connected network be simplified by considering the properties of image recognition?

Why CNN for Image

- Some patterns are much smaller than the whole image



A neuron does not have to see the whole image to discover the pattern.

Connecting to small region with less parameters




Why CNN for Image

- The same patterns appear in different regions.



Why CNN for Image

- The same patterns appear in different regions.



Why CNN for Image

- Subsampling the pixels will not change the object

bird



除非太過頭

bird



subsampling



We can subsample the pixels to make image smaller

→ Less parameters for the network to process the image

Convolutional Neural Network



The whole CNN



The whole CNN



Property 1

- Some patterns are much smaller than the whole image

Property 2

- The same patterns appear in different regions.

Property 3

- Subsampling the pixels will not change the object

Convolution

Max Pooling



Convolution

Max Pooling

Flatten

Can repeat
many times

The whole CNN



Flatten

CNN – Convolution

Those are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1
Matrix

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2
Matrix

⋮

Property 1

Each filter detects a small pattern (3 x 3).

CNN – Convolution


stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1




CNN – Convolution

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0




We set stride=1 below


6 x 6 image

CNN – Convolution


stride=1



6 x 6 image



Filter 1



Property 2

CNN – Convolution

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image


-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

Do the same process for every filter



CNN – Zero Padding



1	-1	-1
-1	1	-1
-1	-1	1


Filter 1

You will get another 6×6 images in this way





Zero padding

CNN – Colorful image



The whole CNN







CNN – Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2



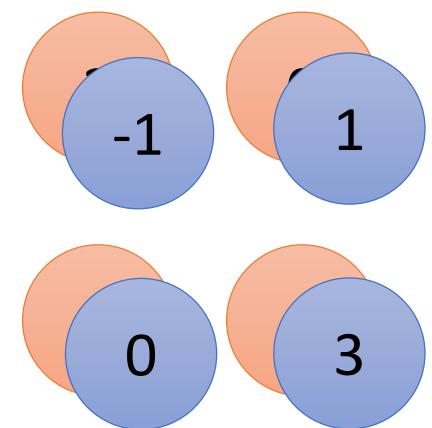
CNN – Max Pooling

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image




New image
but smaller



2 x 2 image


Each filter
is a channel

The whole CNN



The whole CNN

cat dog



Convolution

Max Pooling

A new image



Convolution

Max Pooling

A new image



Flatten


Flatten




Fully Connected
Feedforward network

The whole CNN






(Ignoring the non-linear activation function after the convolution.)






Filter 1



6 x 6 image


Less parameters!







Less parameters!

Even less parameters!





(Ignoring the non-linear activation function after the convolution.)




Dim = $6 \times 6 = 36$

parameters =
 $36 \times 32 = 1152$

Dim = $4 \times 4 \times 2$
 $= 32$


1	0	0	0	0	0	1
0	1	0	0	1	0	0
0	0	1	1	0	0	0
1	0	0	0	0	1	0
0	1	0	0	0	1	0
0	0	1	0	0	1	0

image





convolution


Only $9 \times 2 = 18$
parameters



Max
pooling



Convolutional Neural Network



Learning: Nothing special, just gradient descent

CNN in Keras

Only modified the network structure

```
model.add(Convolution2D(32, 3, 3,  
    border_mode='same',  
    input_shape=(3, 32, 32)))  
model.add(Activation('relu'))
```



```
model.add(Convolution2D(32, 3, 3))  
model.add(Activation('relu'))
```



```
model.add(MaxPooling2D(pool_size=(2, 2)))  
model.add(Dropout(0.25))
```



```
model.add(Convolution2D(64, 3, 3,  
    border_mode='same'))  
model.add(Activation('relu'))
```



```
model.add(Convolution2D(64, 3, 3))  
model.add(Activation('relu'))
```




```
model.add(MaxPooling2D(pool_size=(2, 2)))  
model.add(Dropout(0.25))
```

Code:

https://github.com/fchollet/keras/blob/master/examples/cifar10_cnn.py

```
model.add(Dense(10))
```

```
model.add(Activation('softmax'))
```



```
model.add(Dense(512))
```

```
model.add(Activation('relu'))
```

```
model.add(Dropout(0.5))
```




```
model.add(Flatten())
```





Why CNN for playing Go?

- Some patterns are much smaller than the whole image

Alpha Go uses 5 x 5 for first layer



- The same patterns appear in different regions.



Why CNN for playing Go?

- Subsampling the pixels will not change the object



Max Pooling

How to explain this???

Neural network architecture. The input to the policy network is a $19 \times 19 \times 48$ image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a 23×23 image, then convolves k filters of kernel size 5×5 with stride 1 with the input image and applies a rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a 21×21 image, then convolves k filters of kernel size 3×3 with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size 1×1 with stride 1 with a different bias for each position, and applies a softmax function. The Alpha Go does not use Max Pooling Extended Data Table 3 additionally show the results of training with $k = 128, 256$ and 384 filters.


Variants of Neural Networks

Convolutional Neural
Network (CNN)

Recurrent Neural Network
(RNN) Neural Network with Memory

Example Application

- Slot Filling




Example Application

Solving slot filling by
Feedforward network?

Input: a word

(Each word is represented
as a vector)

Taipei



1-of-N encoding

How to represent each word as a vector?

1-of-N Encoding lexicon = {apple, bag, cat, dog, elephant}

The vector is lexicon size.

$$\text{apple} = [1 \ 0 \ 0 \ 0 \ 0]$$

Each dimension corresponds
to a word in the lexicon

$$\text{bag} = [0 \ 1 \ 0 \ 0 \ 0]$$

The dimension for the word
is 1, and others are 0


$$\text{cat} = [0 \ 0 \ 1 \ 0 \ 0]$$

$$\text{dog} = [0 \ 0 \ 0 \ 1 \ 0]$$


$$\text{elephant} = [0 \ 0 \ 0 \ 0 \ 1]$$

Beyond 1-of-N encoding

Dimension for “Other”



Word hashing



Example Application

Solving slot filling by
Feedforward network?


Input: a word

(Each word is represented
as a vector)


Output:

Probability distribution that
the input word belonging to
the slots


Taipei



Example Application



dest
time of
departure




Neural network
needs memory!

Taipei

Recurrent Neural Network



Recurrent Neural
Network (RNN)




天生的腦



Recurrent Neural Network (RNN)

The output of hidden layer
are stored in the memory.



Memory can be considered
as another input.


RNN

The same network is used again and again.

Probability of
“arrive” in each slot


Probability of
“Taipei” in each slot

Probability of
“on” in each slot



RNN


Prob of “leave”
in each slot



Prob of “Taipei”
in each slot

Different


Prob of “arrive”
in each slot




Prob of “Taipei”
in each slot

The values stored in the memory is different.


Of course it can be deep ...




Bidirectional RNN



Long Short-term Memory (LSTM)







Activation function f is usually a sigmoid function

Between 0 and 1

Mimic open and close gate

$$c' = g(z)f(z_i) + cf(z_f)$$








LSTM

C^{t-1}

vector




LSTM




LSTM

Extension: “peephole”



Multiple-layer LSTM



Don't worry if you cannot understand this.
Keras can handle it.

Keras supports
“LSTM”, “GRU”, “SimpleRNN” layers

This is quite
standard now.




Recurrent Neural Network



天生的腦



Learning Target




Recurrent Neural Network



Learning

Backpropagation
through time (BPTT)




RNN Learning is very difficult in practice.


Unfortunately

- RNN-based network is not always easy to learn

Real experiments on Language modeling



The error surface is rough.



Why?

$$w = 1 \quad \rightarrow \quad y^{1000} = 1$$

$$w = 1.01 \quad \rightarrow \quad y^{1000} \approx 20000$$

$$w = 0.99 \quad \rightarrow \quad y^{1000} \approx 0$$

$$w = 0.01 \quad \rightarrow \quad y^{1000} \approx 0$$


Large
 $\partial L / \partial w$

Small
Learning rate?

small
 $\partial L / \partial w$


Large
Learning rate?

Toy Example



Helpful Techniques

- Advance momentum method
 - Nesterov's Accelerated Gradient (NAG)



Methods:


- Gradient descent
- Momentum
- Nesterov's Accelerated Gradient (NAG)

Source:

<http://www.cs.toronto.edu/~fritz/absps/momentum.pdf>

Helpful Techniques

- Long Short-term Memory (LSTM)
 - Can deal with gradient vanishing (not gradient explode)
 - Memory and input are added
 - The influence never disappears unless forget gate is closed
 - No Gradient vanishing
(If forget gate is opened.)



Helpful Techniques


- Gated Recurrent Unit (GRU)

Simplified LSTM

[Cho, EMNLP'14]


舊的不去、新的不來

GRU has less parameters
than LSTM




Helpful Techniques

Clockwise RNN



[Jan Koutnik, JMLR'14]

Structurally Constrained Recurrent Network (SCRN)



[Tomas Mikolov, ICLR'15]

Vanilla RNN Initialized with Identity matrix + ReLU activation function [Quoc V. Le, arXiv'15]


- Outperform or be comparable with LSTM in 4 different tasks

More Applications

Probability of
“arrive” in each slot

Probability of
“Taipei” in each slot

Probability of
“on” in each slot



Many to one

Keras Example:

https://github.com/fchollet/keras/blob/master/examples/imdb_lstm.py

- Input is a vector sequence, but output is only one vector

Sentiment Analysis

看了這部電影覺
得很高興



Positive (正雷)

這部電影太糟了
.....

Negative (負雷)

這部電影很
棒

Positive (正雷)



Many to Many (Output is shorter)

- Both input and output are both sequences, **but the output is shorter.**
 - E.g. **Speech Recognition**

Problem?

Why can't it be
“好棒棒”


Output: “好棒” (character sequence)



Trimming

好 好 好 棒 棒 棒 棒


Input:



(vector
sequence)


Many to Many (Output is shorter)

- Both input and output are both sequences, **but the output is shorter.**
- Connectionist Temporal Classification (CTC) [Alex Graves, ICML'06][Alex Graves, ICML'14][Hasim Sak, Interspeech'15][Jie Li, Interspeech'15][Andrew Senior, ASRU'15]




Many to Many (No Limitation)

- Both input and output are both sequences *with different lengths*. → *Sequence to sequence learning*
 - E.g. *Machine Translation* (machine learning → 機器學習)



Many to Many (No Limitation)

- Both input and output are both sequences with different lengths. → Sequence to sequence learning
 - E.g. Machine Translation (machine learning→機器學習)




Many to Many (No Limitation)

推	tlkagk:	超	06/12 10:39
推	n:	人	06/12 10:40
推	tion:	正	06/12 10:41
→	host:	大	06/12 10:47
推	:	中	06/12 10:59
推	403:	天	06/12 11:11
推	:	外	06/12 11:13
推	527:	飛	06/12 11:17
→	990b:	仙	06/12 11:32
→	512:	草	06/12 12:15

推 tlkagk: =====斷=====

Many to Many (No Limitation)

- Both input and output are both sequences with different lengths. → Sequence to sequence learning
 - E.g. Machine Translation (machine learning→機器學習)




Add a symbol “==” (斷)

[Ilya Sutskever, NIPS'14][Dzmitry Bahdanau, arXiv'15]


One to Many

- Input an image, but output a sequence of words

[Kelvin Xu, arXiv'15][Li Yao, ICCV'15]



Video Caption Generation



Video frames

Concluding Remarks

Convolutional Neural
Network (CNN)

Recurrent Neural Network
(RNN)

Lecture IV: Next Wave

Outline

Ultra Deep Network

Attention Model

Reinforcement Learning

Realizing what the World Looks Like

Understanding the Meaning of Words

Why Deep?

Ultra Deep Network

Worry about overfitting?

152 layers

Worry about achieving
target first!

This ultra deep network
have special structure.

3.57%

16.4%



AlexNet
(2012)

7.3%




VGG
(2014)

6.7%



GoogleNet
(2014)

Residual Net
(2015)




Ultra Deep Network

- Ultra deep network is the ensemble of many networks with different depth.

Ensemble {
 6 layers
 4 layers
 2 layers

Residual Networks are Exponential
Ensembles of Relatively Shallow Networks
<https://arxiv.org/abs/1605.06431>



Ultra Deep Network

- FractalNet


FractalNet: Ultra-Deep Neural Networks without Residuals
<https://arxiv.org/abs/1605.07648>

Resnet in Resnet

Resnet in Resnet:
Generalizing Residual Architectures
<https://arxiv.org/abs/1603.08029>


Good Initialization?

All you need is a good init
<http://arxiv.org/pdf/1511.06422v7.pdf>




Ultra Deep Network

- Residual Network
- Highway Network




Deep Residual Learning for Image
Recognition
<http://arxiv.org/abs/1512.03385>




Training Very Deep Networks
<https://arxiv.org/pdf/1507.06228v2.pdf>


output layer



output layer



output layer



Highway Network automatically
determines the layers needed!

Input layer

Input layer

Input layer

Outline

Ultra Deep Network

Attention Model


Reinforcement Learning

Realizing what the World Looks Like


Understanding the Meaning of Words

Why Deep?

Attention-based Model




Attention-based Model



Ref:

[http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS_2015_2/Lecture/Attain%20\(v3\).ecm.mp4/index.html](http://speech.ee.ntu.edu.tw/~tlkagk/courses/MLDS_2015_2/Lecture/Attain%20(v3).ecm.mp4/index.html)

Reading Comprehension



Reading Comprehension

- End-To-End Memory Networks. S. Sukhbaatar, A. Szlam, J. Weston, R. Fergus. NIPS, 2015.


The position of reading head:

Story (16: basic induction)	Support	Hop 1	Hop 2	Hop 3
Brian is a frog.	yes	0.00	0.98	0.00
Lily is gray.		0.07	0.00	0.00
Brian is yellow.	yes	0.07	0.00	1.00
Julius is green.		0.06	0.00	0.00
Greg is a frog.	yes	0.76	0.02	0.00
What color is Greg? Answer: yellow		Prediction: yellow		

Keras has example:


https://github.com/fchollet/keras/blob/master/examples/babi_memnn.py

Visual Question Answering



source: <http://visualqa.org/>

Visual Question Answering



Visual Question Answering



- Huijuan Xu, Kate Saenko. Ask, Attend and Answer: Exploring Question-Guided Spatial Attention for Visual Question Answering. arXiv Pre-Print, 2015

Is there a red square on the bottom of the cat?

GT: yes




Prediction: yes




Video Caption Generation

Memory: video frames
Output: video description



Video Caption Generation

Memory: video frames
Output: video description



Video Caption Generation

- Demo: 曾柏翔、盧宏宗、吳柏瑜

Outline

Ultra Deep Network

Attention Model

Reinforcement Learning


Realizing what the World Looks Like

Understanding the Meaning of Words

Why Deep?


Only Supervised Learning until now

- Network is a function. In supervised learning, the input-output pair is given in the training data




Supervised v.s. Reinforcement

- Example: Dialogue Agent for 訂票系統



Supervised v.s. Reinforcement

- Example: Dialogue Agent for 訂票系統




Supervised v.s. Reinforcement

- Supervised



You have to “greeting”

- Reinforcement



Bad

Supervised v.s. Reinforcement

- Playing GO
 - Supervised: 看著棋譜學



下一步：
“5-5”



下一步：
“3-3”

- Reinforcement Learning

初手天元 →下了好幾百手 → Win!

Alpha Go is supervised learning + reinforcement learning.

To learn deep reinforcement learning

- Lectures of David Silver
 - <http://www0.cs.ucl.ac.uk/staff/D.Silver/web/Taching.html>
 - 10 堂課 (1:30 each)
- Deep Reinforcement Learning
 - http://videolectures.net/rldm2015_silver_reinforcement_learning/

Outline

Ultra Deep Network

Attention Model

Reinforcement Learning


Realizing what the World Looks Like

Understanding the Meaning of Words


Why Deep?

Does machine know what the world look like?

Ref: <https://openai.com/blog/generative-models/>



Draw something!



Deep Dream


- Given a photo, machine adds what it sees



<http://deepdreamgenerator.com/>

Deep Dream

- Given a photo, machine adds what it sees



<http://deepdreamgenerator.com/>

Deep Style


- Given a photo, make its style like famous paintings



<https://dreamscopeapp.com/>


Deep Style

- Given a photo, make its style like famous paintings




<https://dreamscopeapp.com/>

Deep Style




Generating Images (無中生有)

- Training a decoder to generate images is **unsupervised**





Training data is a lot of images




Auto-encoder

Not state-of-the-art approach





Generating Images

- Training a decoder to generate images is **unsupervised**
- Variation Auto-encoder (VAE)
 - Ref: **Auto-Encoding Variational Bayes**,
<https://arxiv.org/abs/1312.6114>
- Generative Adversarial Network (GAN)
 - Ref: **Generative Adversarial Networks**,
<http://arxiv.org/abs/1406.2661>



Which one is machine-generated?




Ref: <https://openai.com/blog/generative-models/>


Generating Images by RNN




color of
2nd pixel



color of
3rd pixel



color of
4th pixel




color of
1st pixel

color of
2nd pixel

color of
3rd pixel

Generating Images by RNN

- **Pixel Recurrent Neural Networks**
 - <https://arxiv.org/abs/1601.06759>



Outline

Ultra Deep Network

Attention Model

Reinforcement Learning

Realizing what the World Looks Like

Understanding the Meaning of Words

Why Deep?

Machine Reading


- Machine learn the meaning of words from reading a lot of documents without supervision



<http://top-breaking-news.com/>


Machine Reading

- Machine learn the meaning of words from reading a lot of documents without supervision



Machine Reading

- Generating Word Vector/Embedding is **unsupervised**



Training data is a lot of text



Machine Reading

- Machine learn the meaning of words from reading a lot of documents without supervision
- A word can be understood by its context

蔡英文、馬英九 are something very similar



You shall know a word by the company it keeps

馬英九 520宣誓就職

蔡英文 520宣誓就職



Word Vector



Source: <http://www.slideshare.net/hustwj/cikm-keynotenov2014>

Word Vector

$$\approx V(Berlin) - V(Rome) + V(Italy)$$

- Characteristics

$$V(hotter) - V(hot) \approx V(bigger) - V(big)$$

$$V(Rome) - V(Italy) \approx V(Berlin) - V(Germany)$$

$$V(king) - V(queen) \approx V(uncle) - V(aunt)$$

- Solving analogies



$$Rome : Italy = Berlin : ?$$

Compute $V(Berlin) - V(Rome) + V(Italy)$

Find the word w with the closest $V(w)$

Machine Reading

- Machine learn the meaning of words from reading a lot of documents without supervision



Machine learns to understand 鄉民用語 via reading the posts on PTT

Demo

- Model used in demo is provided by 陳仰德
- Part of the project done by 陳仰德、林資偉
- TA: 劉元銘
- Training data is from PTT (collected by 葉青峰)

Outline

Ultra Deep Network

Attention Model

Reinforcement Learning

Realizing what the World Looks Like

Understanding the Meaning of Words

Why Deep?

Deeper is Better?

Layer X Size	Word Error Rate (%)
1 X 2k	24.2
2 X 2k	20.4
3 X 2k	18.4
4 X 2k	17.8
5 X 2k	17.2
7 X 2k	17.1

Not surprised, more parameters, better performance

Seide, Frank, Gang Li, and Dong Yu. "Conversational Speech Transcription Using Context-Dependent Deep Neural Networks." *Interspeech*. 2011.


Universality Theorem

Any continuous function f

$$f : R^N \rightarrow R^M$$

Can be realized by a network
with one hidden layer


(given **enough** hidden
neurons)



Reference for the reason:
<http://neuralnetworksanddeeplearning.com/chap4.html>

Why “Deep” neural network not “Fat” neural network?

Fat + Short v.s. Thin + Tall




Fat + Short v.s. Thin + Tall

Layer X Size	Word Error Rate (%)	Layer X Size	Word Error Rate (%)
1 X 2k	24.2		
2 X 2k	20.4		
3 X 2k	18.4		
4 X 2k	17.8		
5 X 2k	17.2	1 X 3772	22.5
7 X 2k	17.1	1 X 4634	22.6
		1 X 16k	22.1

Why?

Seide, Frank, Gang Li, and Dong Yu. "Conversational Speech Transcription Using Context-Dependent Deep Neural Networks." *Interspeech*. 2011.

Analogy



Logic circuits

- Logic circuits consists of **gates**
- **A two layers of logic gates** can represent **any Boolean function.**
- Using multiple layers of logic gates to build some functions are much simpler



less gates needed



- Neural network consists of **neurons**
- **A hidden layer network** can represent **any continuous function.**
- Using multiple layers of neurons to represent some functions are much simpler



less parameters




less data?

This page is for EE background.

Modularization


- Deep → Modularization





Modularization

Each basic classifier can have sufficient training examples.

- Deep → Modularization



Classifiers for the attributes




v.s.

v.s.

Modularization


can be trained by little data

- Deep → Modularization



Modularization

- Deep \rightarrow Modularization \rightarrow Less training data?



The most basic
classifiers


Use 1st layer as module
to build classifiers

Use 2nd layer as
module

Modularization

Reference: Zeiler, M. D., & Fergus, R. (2014). Visualizing and understanding convolutional networks. In *Computer Vision–ECCV 2014* (pp. 818-833)

- Deep → Modularization




The most basic
classifiers

Use 1st layer as module
to build classifiers


Use 2nd layer as
module

SVM




Source of image: http://www.gipsa-lab.grenoble-inp.fr/transfert/seminaire/455_Kadri2013Gipsa-lab.pdf


Deep Learning



Boosting



Deep Learning



More Reasons

- Do Deep Nets Really Need To Be Deep? (by Rich Caruana)
- <http://research.microsoft.com/apps/video/default.aspx?id=232373&r=1>

Do deep nets really
need to be deep?

Rich Caruana
Microsoft Research

Lei Jimmy Ba
MSR Intern, University of Toronto

Thanks also to: Gregor Urban, Krzysztof Geras, Samira Kahou, Abdelrahman Mohamed, Jinyu Li, Rui Zhao, Jui-Ting Huang, and Yifan Gong

Yes!

Thank You

Any Questions?

Concluding Remarks

Today's Lecture

Lecture I: Introduction of Deep Learning



Lecture II: Tips for Training Deep Neural Network



Lecture III: Variants of Neural Network



Lecture IV: Next Wave

Some Opinions

- Also learn other machine learning methods


不想知道 Deep Learning 以外的方法



吾名乃惠惠

Some Opinions

- In some situations, the simpler machine learning methods can be very powerful.



Some Opinions

<http://www.baike.com/gwiki/%E5%AF%92%E6%AD%A6%E7%BA%AA%E5%A4%A7%E7%88%86%E5%8F%91>

- 寒武纪大爆炸



已經有一些生物滅絕了

Some Opinions

- Deep Learning is still at the phase of “神農嘗百草”
- Lots of questions still do not have answers
- However, probably also easy to enter



http://orchid.shu.edu.tw/upload/article/20110927181605_1_pic.png

如果你想“深度學習 深度學習”

- “Neural Networks and Deep Learning”
 - written by Michael Nielsen
 - <http://neuralnetworksanddeeplearning.com/>
- “Deep Learning”
 - Written by Yoshua Bengio, Ian J. Goodfellow and Aaron Courville
 - <http://www.iro.umontreal.ca/~bengioy/dlbook/>
- Course: Machine learning and having it deep and structured
 - http://speech.ee.ntu.edu.tw/~tlkagk/courses_MLSD15_2.html

給資料科學愛好者

- 台大電機系於台大電信所成立「資料科學與智慧網路組」，開始招收碩、博士生
- 今年秋天開始報名

Kibana

Machine Learning

Elasticsearch Deep Learning

NTU GICE Data Scientist Logstash
Machine Learning Python

Data Science and Deep Learning Kibana

Smart Networking Logstash

Data Scientist Deep Learning
Machine Learning

Python